

Latent Representation of Topographic Classes in Remote Sensing Image Data using Autoencoders

Hannes Stärk

August 19, 2019

Contents

1	Vorwort	2
2	Kurzfassung	2
3	Abstract	2
4	Introduction	2
4.1	Topic Overview	2
4.2	Research Questions	2
5	Background	2
5.1	Remote Sensing Imagery	2
5.2	Artificial Neuron	3
5.3	Artificial Neural Network	3
5.4	Convolutional Neural Networks	5
5.5	Autoencoders	7
5.6	Kullback-Leibler Divergence	8
5.6.1	Entropy	8
5.6.2	Cross Entropy	9
5.6.3	Kullback-Leibler Divergence	9
5.7	Variational Autoencoders	10
5.8	t-Distributed Stochastic Neighbor Embedding	11
6	Related Work	11
6.1	Convolutional Network Architecture	11
6.2	Variational Autoencoders and Generative Networks	12
7	Methodology	12
7.1	Environment	12
7.1.1	Hardware	12
7.1.2	Software	13

7.2	Datasets	13
7.3	Architecture	14
7.3.1	Pure Convolutional Architecture	15
7.3.2	Max Pooling Architecture	16
7.4	Understanding Latent Space	16
8	Experiments	17
8.1	Variational Autoencoder Architecture	17
8.2	Latent Space	17
9	Conclusion and Future Work	17
9.1	Conclusion	17
9.2	Future Work	17

1 Vorwort

2 Kurzfassung

3 Abstract

4 Introduction

4.1 Topic Overview

4.2 Research Questions

5 Background

This section contains basic knowledge that is necessary for the understanding of the rest of the thesis. It briefly touches on artificial neural networks and convolutional neural networks in the beginning. The subsection after that goes on to describe autoencoders in general. Next up, Kullback-Leibler divergence is explained in detail as it is an essential part for the variational autoencoders as they are used in this work. Kullback-Leibler divergence also plays an important role in t-distributed stochastic neighbor embedding. The difference between variational autoencoders and the standard undercomplete autoencoder is covered in the next subsection. At last, the dimensionality reduction method t-distributed stochastic neighbor embedding is brought up since it is used in this work to gain an understanding of the high dimensional latent space of the autoencoders.

5.1 Remote Sensing Imagery

Logic: " [...]"

5.2 Artificial Neuron

Logic: "Inspired by biological neurons an artificial neuron has one or more [...] "The sentence has to be split into two different statements: Firstly, artificial neurons are inspired by [...]. Secondly, in computer science the artificial neurons are seen as functions with input parameter [...]. The following brief explanation of the function is good.

The output models the **Prior knowledge:** "axon" The reader has certainly to browse the terms axon and dendrites?! That complies not with the demand of continuously reading. Do you need really these terms without a graphic? of a biological neuron to which other biological neurons can connect via their dendrites which is similar to an artificial neuron using the output of another **Repetition:** "artificial neuron" artificial neuron as its input.

By adjusting the separate weights of the inputs a single neuron is able to model simple functions like the logic *and* or the logic *or* but it is not abled to solve other trivial problems like the exclusive or (XOR).

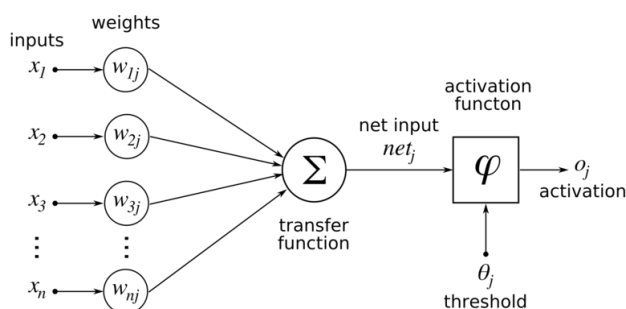


Figure 1: An artificial neuron. Image taken from (Chrislb 2005)

5.3 Artificial Neural Network

Global: Paragraph Please adapt the order of the following paragraph to the last: Firstly, the functional model can be written as $f(x) = y$. Moreover the weights and the bias has to be introduced in a corresponding function. After that you can present the idea of layers.

Logic: "By taking layers of multiple artificial neurons [...]" In the last paragraph you wrote about single neurons, then we need a halfsentence that a layer consists of multiple [...]. A verb is missing here? next layer, functions $f(x) = y$ with input and output vectors x and y of any size can be modeled.

Prior Knowledge: "The layers between the input and output layer are called hidden layers." Before introducing hidden lyers refer to the multilayer perceptron! A MLP consists of at least three layers of nodes. The term "at least" is highly informative Arbitrarily complex functions **Prior Knowledge:** "could" "can" or "may", it is not as hypothetical as it sounds be modeled by **Prior Knowledge:** "three layers" only three layers? **Grammar:** "if they mays" could, if, may: that sounds totally

hypothetical. But besides that, I would not combine if and may in that way. contain arbitrarily many neurons **Sentence:** but in practice most of the time deep neural networks with two or more hidden layers are used (Géron 2017). **Breaking sentence into two parts.** If all the neurons of a previous layer are connected to each neuron in the following layer the layers are called fully connected which is the case for every layer in the artificial neural network depicted in Figure 2.




Figure 2: An artificial neural network with three fully connected layers. Image taken from (Glosser.ca 2013)

Figure: Please refer in 5.3 slightly more to that graphic as well as to that of 5.2

Notation: In 5.3 you used function f . Is function n a fundamentally new model? Moreover you did not introduced w in your formula. In future you might work with highly complex functions, which have to be exactly initiated.

The process of learning an artificial neural networks $n(w, x) = z$, **Order:** "with weights w " introduce weights in 5.3,

Order: "refers to the process of iteratively adjusting w in such a way that $n(w, x)$ more closely resembles the desired function that should be learned $f(x)$." Firstly your formula describes not an iterative process, you have to build a nested function to achieve an iterative process. Secondly, I don't know what you want to express with "desired function", might it be the desired function (output) of the function "model". I think, the term "function" can be misunderstood.

Reference: "This can be achieved" What can be achieved? by taking possible inputs x_i for which the desired outputs

Notation: \hat{z}_i in 5.3 you introduced y as output. Furthermore you present the ideal of ground truth, or rather ideal data but only in your formula are known and optimizing a **Prior Knowledge** "loss function What is a loss function $l(n(w, x_i), \hat{z}_i)$ for w **Logic:** "and repeating this process with each pair of inputs x_i with desired outputs \hat{z}_i ." Firstly, why have I to repeat that process. That is a slightly complicated way to introduce the concept of learning without the use of the terms "training" or "learning". Secondly, the ground truth data is not desired, they are merely given. I think, you want to say the desired output are very similar. Thirdly, the output is commonly not presented by the hat? . The set of pairs of those inputs and desired outputs is called the training set. Intuitively the optimized loss function

can be seen as the **Term:** "goal" "objective"? of how the output of the artificial neural network should resemble **Term:** "the known desired outputs." "ground truth?" **Repetition:** "desired"

Prior Knowledge: "The optimization is most commonly done via gradient descent or adjusted versions of gradient descent." The objective (is there a representable general objective?) can be seen as an optimization problem. What is the gradient descent for? This involves calculating the partial derivatives for each component of w . For this reason the activation functions of the artificial neurons and the loss function have to be differentiable. More detailed information about gradient descent can be read in Deep Learning (Goodfellow et al. 2016) and information about the commonly used optimizers that improve gradient descent can be found in Hands-On Machine Learning (Géron 2017).Nice!

5.4 Convolutional Neural Networks

Global: You need an example. Please use any example, show a simple picture and introduce a function with simple objective without formula but referring to that of the last subsection. This subsection suffers from sparsely describing figurative. Hence, there are several possible misconceptions.

Global:The next two paragraphs might briefly introduce the concept of convolutional neural nets, but I can still not really understand why it brings a clear added value. You have to explain it. Additionally the specification of a potentially use of convolutional neural nets involves architecture issues (kernels, padding). You are introducing such concepts but it is rather not clear why you do this.

When working with neural networks that should model a function with data like images as inputs, **Term:** "most of the time." "mostly" of the time convolutional neural networks (CNNs) are used that contain convolutional layers. A convolutional layer is not connected to every neuron of the previous layer but instead only to small rectangle of neurons of the previous layer. For images this means that in the first convolutional layer the neurons do not take every pixel of the image as input but instead only a small piece of the image. This way a hierarchical structure is created where the first hidden layer learns lower level features and the following layers learn **Term:** "ever" ?? higher level features.

Term: Also by only connecting a neuron to a few neurons of the previous layer the number of weights is far lower which means there is much less memory and time needed for training the network. In opposite to [...]

Sentences Is the next paragraph to be seen as a kind of brainstorming? Your wealth of expressions is massively reduced?

Sentence: "The rectangle a neuron in a convolutional layer takes its inputs from is called kernel" That is not really a sentence...?? and has a kernel size specifying the height and width of the kernel. If the kernel of a neuron laps over the edge of the input space different kinds of padding can be used for the values of the kernel where there are no input values. These kernels, also called filters, are basically maps of weights that are adjusted while training which way they learn features of the input. In practice multiple neurons have the same input space

leading to the ability to extract multiple features that can be processed further in the following layer. The amount of neurons with the same input space is referred to as the number of filters since each one of those neurons has its own filter with its respective weights and all those filters are applied to the same piece of the input. This is visualized in Figure 3.




Figure 3: A convolutional network with 12 filters in the first convolutional layer and 7 filters in the second convolutional layer. Each filters can learn different features and an RGB image is depicted as the input here. Image taken from (Géron 2017)

Instead of having a kernel in each possible area of the input there can also be a distance between each kernel which is called stride. This can be used to decrease the dimensionality of the output of a layer. For example with strides of 2 the output dimensions are halve that of the input dimensions.

I really appreciate your work!! From an autistic point of view there is a lot to do, but at the moment the aspiration is somewhat higher than required. You can include the questions and comments into your work, but you can also focus on programming. Good luck.

Pooling A pooling layer replaces the value of an output with a summary of the nearby outputs. For instance max pooling only takes the maximum output of all outputs in a surrounding rectangle. Note that the pooling layers just perform an operation on multiple outputs of a previous layer and therefore do not have weights. The use of pooling layers increases the networks tolerance to small changes in the inputs (Goodfellow et al. 2016) which is most of the time desired since the network should learn features instead of exact pixel locations. Many popular CNN architectures like AlexNet (Krizhevsky et al. 2012) use pooling layers between the convolutional layers. In these CNNs the pooling layers also

often are used to reduce the dimension of the layers.

5.5 Autoencoders

An autoencoder is an artificial neural network that aims to reproduce its input. The first part of an autoencoder is the Encoder $f(x) = z$ that takes an input x and maps it to a latent code z . The second part is the decoder $g(z) = x'$ that tries to generate a reproduction x' similar to the input x that the latent code z was produced from. Since the output of the neural network should be the same as the input there is no need for labeled data and the network can be learned unsupervised. Just having a network that is the identity function would be useless but by constraining the autoencoder in specific ways it can be forced to learn useful traits leading to possibilities in pretraining networks or randomly generating content similar to the content the autoencoder was trained with. One common constraint for autoencoders is choosing a dimension for the latent code that is smaller than the dimension of the input. This means that the encoder is forced to learn to extract only the most important features from the input and the latent code is a representation of those most important features in a lower dimensionality than the input. Such an autoencoder, called undercomplete autoencoder, is visualized in Figure 4.




Figure 4: An undercomplete autoencoder has a latent code z whose dimension is smaller than the dimension of the input x . Image taken from (Chervinskii 2015)

Undercomplete autoencoders can be used for pretraining by taking the trained Encoder that has learned the most important features of the input and using it for the actual desired task like classification. Another obvious application would be dimensionality reduction.

5.6 Kullback-Leibler Divergence

5.6.1 Entropy

A part of the loss function used in the variational autoencoder is based on Kullback-Leibler divergence. To understand Kullback-Leibler divergence it seems necessary to explain entropy from the field of information theory. In short, entropy is a measure for the minimum average size an encoding for a piece of information can possibly have.

Suppose there is a system $S1$ that can have four different states a, b, c, d and every one of those states is equally likely to occur, that means the probability $P(x)$ of each state x is $1/4$. Now the goal is to losslessly transmit all information about that system with the minimum average amount of bits. That can be done with only two bits for example like this

$$a : 00 \quad b : 01 \quad c : 10 \quad d : 11$$

However, if $P(a) = 1$ and $P(b) = P(c) = P(d) = 0$, zero bits will suffice to encode the information since it is always certain that the system is in state a . So the entropy of the system clearly depends on the probabilities of each state. To see in which way, one can consider the system $S2$ with $P(a) = 1/2$, $P(b) = 1/4$, $P(c) = 1/8$ and $P(d) = 1/8$. In that case it would be best to encode the state with the highest probability with as few bits as possible since it has to be transmitted the most often. That means a is encoded with one bit as 0. When decoding the information there must be no ambiguities so while the encoding for b has to start with a 1 it cannot be 1 since we need to encode two more states so $b : 10$. Additionally if $c : 11$ there would be no space left for d : say $d : 111$ then if the transmitted information is $111111\dots$ it could either be decoded to $ccc\dots$ or $dd\dots$. So c should rather be encoded as 110 which way $d : 111$ works. In the end a valid encoding that can transmit all information with the minimum average amount of bits is

$$a : 0 \quad b : 10 \quad c : 110 \quad d : 111$$

Here the states c, d are encoded with three bits instead of the two bits in the first example. But c and d are transmitted far less often than a which now only needs one bit. To be more precise half of all transmissions have one bit. Additionally a quarter of all transmissions have two bits. The sum of those probabilities multiplied with the respective amount of bits is the average amount of bits needed to transfer the information in a given encoding. So in the example, with $f(x)$ as the number of bits that encode a state x , that turns out to be $P(a)f(a) + P(b)f(b) + P(c)f(c) + P(d)f(d) = 1.75$. That means for $S2$ on average you only need 1.75 bits to encode a state and since that is also the minimum 1.75 is the entropy of $S2$.

In general in an optimal encoding $f(x)$ is the same as $\log_2 \frac{1}{P(x)}$. For $S1$ $P(x)$ is $1/4$ so the number of bits for x is $\log_2(4) = 2$ what matches the two bits the first encoding uses for the states of $S1$.

The entropy H of a system with a set of discrete events X and the probability distribution $P(x)$ for each $x \in X$ is

$$H(P) = \sum_{x \in X} P(x) \log_2 \frac{1}{P(x)} = - \sum_{x \in X} P(x) \log_2 P(x) \quad (1)$$

This is often written as the expectation for a given state x under the distribution P .

$$H(P) = E_{x \sim P}[-\log_2 P(x)]$$

Intuitively if a system has high entropy, the size of the encodings are high on average and many states have small probabilities. This means it is hard to predict what state the system will be in at a given time since there is no state that can be guessed with high confidence. If entropy is low, zero for example, one can be confident that the system is in a certain state like in the previous example with $P(a) = 1$.

5.6.2 Cross Entropy

If the real distribution P of a system is unknown an estimate distribution Q could be guessed and encoding sizes $-\log_2 Q(x)$ can be produced which will not be optimal for the true distribution P . Now with some data gathered and P known the used encoding sizes can be cross-checked with the expectation under the actual distribution resulting in the cross entropy $H(P, Q)$

$$H(P, Q) = E_{x \sim P}[-\log_2 Q(x)] = - \sum_{x \in X} P(x) \log_2(Q(x)) \quad (2)$$

In machine learning tasks regarding classification this is often used as a loss function since the label of a piece of data gives us a distribution P with absolute certainty and $H(P) = 0$. With the inaccurate distribution Q that the model estimates $H(P, Q)$ will be greater than zero unless $P = Q$ where $H(P, Q) = H(P, P) = H(P) = 0$. So the learning algorithm can try to minimize $H(P, Q)$.

5.6.3 Kullback-Leibler Divergence

Having computed the entropy $H(P)$ and cross-entropy $H(P, Q)$ of two distributions P, Q it is possible to compare those distributions by comparing $H(P)$ and $H(P, Q)$ through subtraction

$$\begin{aligned} D_{KL}(P \parallel Q) &= H(P, Q) - H(P) \\ &= E_{x \sim P}[-\log_2 Q(x)] - E_{x \sim P}[-\log_2 P(x)] \\ &= E_{x \sim P}[-\log_2 Q(x) + \log_2 P(x)] \\ &= E_{x \sim P}[\log_2 \frac{P(x)}{Q(x)}] \\ &= \sum_{x \in X} P(x) \log_2 \frac{P(x)}{Q(x)} \end{aligned} \quad (3)$$

where $D_{KL}(P \parallel Q)$ is called the Kullback-Leibler divergence of P and Q . This works because $D_{KL}(P \parallel Q)$ is zero if Q and P are the same since that means $H(P, Q) = H(P)$. On the opposite if Q is different from P then $H(P, Q)$ is greater than $H(P)$ and therefore the KL divergence is greater than zero proportional to how different Q and P are. In summary Kullback-Leibler divergence is a measure of how different two probability distributions are that is zero if they are the same and greater than zero if not.

5.7 Variational Autoencoders

Standard autoencoders are fairly limited in their generative capabilities or in their capabilities to interpolate in the latent space and the problem lies in their latent code. For generation tasks a random latent code z is sampled from the space of possible latent codes. But with a standard autoencoder this space of possible latent codes is most likely not continuous and chances are high that the sampled z is intuitively speaking from a gap in the latent space that the decoder has not been trained for. Therefore it will not produce realistic outputs. This discontinuous latent space also means that interpolation in it does not work well.

To solve this in a variational autoencoder the encoder does not produce discrete values but rather probability distributions p_i by generating two vectors of the same size as z where one vector μ contains means and the other vector σ contains standard deviations. Then each z_i is determined by sampling it from $p_i = N(\mu_i, \sigma_i)$. As a result the decoder is forced to not only just learn how to decode single, discrete latent codes but also the possible variations introduced by the random sampling. This leads to the desired continuous decodable latent space. However, the encoder could still produce σ sigma that are very small and μ that are far apart. To prevent this the encoder is forced to produce distributions that are similar to a normal distribution with standard deviation 1 and mean 0 by adding a second loss function next to the loss $L(x, x')$ between the input and output.

As discussed in 5.6 Kullback-Leibler divergence can be seen as a measure for the difference of two probability distributions. This means that it can be used as the second loss function for the variational autoencoder that should now also minimize the sum of all the Kullback-Leibler divergences between the output distribution and $N(0, 1)$.

To achieve the goal of generating new content that is similar to the input that the variational autoencoder was trained with, one can sample each component for a latent code from $N(0, 1)$ and use it as the input for the decoder. Also smooth interpolation in the latent space is possible leading to possibilities for combining content in interesting ways.

Further details about variational autoencoders can be found here (Doersch 2016).

5.8 t-Distributed Stochastic Neighbor Embedding

The method t-distributed stochastic neighbor embedding (t-SNE) is a machine learning technique for reducing dimensionality. It was first introduced by Laurens van der Maaten and Geoffrey Hinton (Maaten et al. 2008) who claim that t-SNE would be especially well suited for creating single visualizations that reveal the structure of the high dimensional data. Their paper is also where most of the information for this subsection is drawn from.

The method generates an initial embedding in the low dimensional space by sampling from a normal distribution centered around the origin. The learning process improves this embedding by first calculating a representation of the similarities between all points in the high dimensional space. For the similarity of point x_i to x_j this representation is a conditional probability $p_{ij} = \frac{r_{ij} + r_{ji}}{2N}$, so the conditionals are symmetric, with r_{ij} given by the following.

$$r_{ij} = \frac{\exp(-|x_i - x_j|^2 / 2\sigma_i^2)}{\sum_k \sum_{l \neq k} \exp(-|x_k - x_l|^2 / 2\sigma_i^2)} \quad (4)$$

For the similarity of points y_i and y_j from the low dimensional space the representation is a conditional probability q_{ij} given by the following.

$$q_{ij} = \frac{(1 + -|x_i - x_j|^2)^{-1}}{\sum_k \sum_{l \neq k} (1 + -|x_i - x_j|^2)^{-1}} \quad (5)$$

Then the mistake between the representation in high dimensional space and in low dimensional space is given by the Kullback-Leibler divergence 5.6 of high dimensional and low dimensional conditionals which is $\sum_i \sum_{j \neq i} p_{ij} \log \frac{p_{ij}}{q_{ij}}$.

The Kullback-Leibler divergence is then minimized with gradient descent, moving the points in the embedding to a more optimal position.

6 Related Work

This section presents concrete insights from related literature that influenced the used methodology. It is divided into two subsections. The first one, being convolutional network architecture, covers common practice about hyper parameters and other choices for convolutional models while the second subsection focuses on generative aspects and autoencoders.

6.1 Convolutional Network Architecture

In their paper *Striving for simplicity: The All Convolutional Net* (Springenberg et al. 2015) the authors find that including pooling layers such as max pooling into a convolutional architecture does not necessarily mean an improvement. When using a large enough network the invariances that should be introduced by pooling layers can also be learned in a network with convolutional layers only. Thus, in certain cases, convolutional layers with strides of two can be used to reduce dimensionality in a CNN instead of pooling layers which is why

both options are explored in this work.

The paper *Systematic evaluation of CNN advances on the ImageNet* (Mishkin et al. 2016) analyzes the impact of multiple learning parameter and architecture choices on the performance of convolutional neural networks. One conclusion it reaches is the suggestion to use batch sizes around 128 or 256.

6.2 Variational Autoencoders and Generative Networks

In *A Note on the Evaluation of Generative Models* (Theis et al. 2015) it is made clear that there are many difficulties in measuring the optimization and performance of generative models. A clear conclusion is that there is no one fits all metric for the quality of models like variational autoencoders. While the intuitive evaluation of visual results favors overfitted models, pure reliance on numerical measures is not appropriate when the goal is image synthesis. Overall the quality of evaluation measure is largely dependent on the application and goal.

7 Methodology

In this section the subsection environment will explain what tools were used to program the autoencoders and what hardware the developed models were trained and written on. The subsection Datasets contains information about the remote sensing image data that was analyzed using the autoencoders. The architecture subsection lays out the design of the different models regarding the layers and loss functions. In the last subsection latent space there are the techniques used to visualize and understand the latent code between the encoder and decoder with comparisons of the latent space between normal autoencoders and variational autoencoders.

7.1 Environment

7.1.1 Hardware

Training autoencoders with larger images of multiple channels takes a lot of computation especially if the features that should be learned are as complicated and abstract as the topography of remote sensing data. As this is very time consuming the training of the models was mainly done on a remote machine from the Leibniz University in Hannover that could run 24/7. Only the programming of the models and tests with small datasets took place on a local personal computer that had no GPU that was capable of training models in the given scenario. The following are the specifications of the used systems:

Personal Computer

AMD Ryzen 7 1700, 16 GB RAM, NVIDIA GeForce GTX 760 2GB

Remote Machine

Lenovo Legion Y520T-25IKL, Intel i7-7700, 8GB RAM, NVIDIA GeForce GTX 1060 3GB

7.1.2 Software

The programming language used for the work is Python (Rossum 1995) with the development environment PyCharm locally and Jupyter Notebooks on a remote machine. The machine learning framework Tensorflow was used which is developed by Google and offers cross-platform support, running on most CPUs and GPUs (Martín Abadi et al. 2015). Whenever possible Tensorflows implementation of the Keras API specification was used which is a high level API for training machine learning models. For dimensionality reduction of the latent space, Scikit-Learn was used (Pedregosa et al. 2011).

7.2 Datasets

The available data are 1024x1024 images of the two United States cities Jacksonville in Florida and Omaha in Nebraska taken from the US3D Dataset that was partially published to provide research data for the problem of 3D reconstruction (Bosch et al. 2019). The images for each recorded area cover one square kilometer and can be divided into four categories with the first one being multispectral satellite images with eight channels (MSI). From the MSI data three channels were extracted and used as red, green and blue intensities (RGB). Thirdly there are digital surface models (DSM) and Lastly semantic labeling with five different categories.

The MSI data was collected by the WorldView-3 satellite of Digital Globe from 2014 to 2016. Thereby the images were taken in different seasons and times of day leading to great differences in their appearance regarding for instance shadows, reflections, overall brightness or clouds. This is an advantage for training models that are capable of processing data with similar differences in appearance. In the MSI dataset a single picture consists of eight channels for eight different bands of the spectrum with a ground sample distance of 1.3 meters. The eight channels of the imagery correspond to the following wavelengths:

- Coastal: 400 - 450 nm
 - Blue: 450 - 510 nm
 - Green: 510 - 580 nm
 - Yellow: 585 - 625 nm
- 1. Red: 630 - 690 nm
 - 2. Red Edge: 705 - 745 nm
 - 3. Near-IR1: 770 - 895 nm
 - 4. Near-IR2: 860 - 1040 nm

Three of those channels were extracted and used as RGB data. Each pixel of an image is described by three bytes representing the intensity of the wavelengths of the reflected light corresponding to either red, green or blue.

The DSM data was collected using light detection and ranging technology (Lidar) which measures the distance to points of earths surface. This distance is proportional to the value of the single channel that each pixel of the DSM has. Lastly there are semantic labeled pictures with one channel of a single byte that encodes one of five different topographic classes. Those classes are vegetation, water, ground, building and clutter. The semantic labeling was done automatically from lidar data but manually checked and corrected afterwards. Those four categories of data all cover one square kilometer in each image. Additionally they contain a lot of oblique view on buildings and other valuable features like clear shadows making the data ideal for training models that should detect those features.

In the dataset there are 2,783 1024×1024 RGB images available. Since the autoencoder should be able to distinguish between categories like shadows and vegetation each image is split into 64 128×128 pictures resulting in 178,112 total training images. Experimentation with larger image dimensions never produced reconstructions of acceptable quality regarding pure visual inspection. Another option would be to resize larger sections of the original images to 128×128 . However, smaller picture sections are more likely to have a dominant feature like containing mainly shadows, buildings or vegetation. This is desirable since the variational autoencoder should learn to distinguish those features and to cluster them in an unsupervised manner.

7.3 Architecture

The first general structure of the vanilla autoencoder and the variational autoencoder is a combination of the convolutional autoencoder and the variational autoencoder presented in the second edition of *Hans-On Machine Learning* (Géron 2017). That model is adjusted to work with 128×128 RGB images and various different architectural changes are tested. These experiments are first conducted with only 3000 training images and 200 validation images since the training should be fast to allow testing many different architectures. The best performing models are then evaluated and compared after training with all 178,122 pictures. Here the performance is measured regarding the quality of the reconstruction and not the quality of the latent space. In all these test the dimension of the latent code is 1024 while the architectures allow to alter the

coding size for the experiments regarding the latent space 8.2. The described experiments regarding the architecture and their evaluation can be found in section 8.1.

The results were that a purely convolutional network without pooling performed the best. However, to inspect the effects of pooling on the latent code an architecture with pooling is used as the second model in the latent space experiments. This model yields results that are only a little worse than the ones from the purely convolutional network but the training times were almost twice as long.

The following are choices that apply to both architectures. The RGB input values between 0 and 255 are first normalized. Each value is divided by 255 which way there are only floats between 0 and 1. The last deconvolutional layer in the decoder network (Figure 6) uses the sigmoid activation function $S(x) = \frac{e^x}{e^x + 1}$. This guarantees that the values in the output image are all between 0 and 1 and can be interpreted as RGB values. Every other layer uses a scaled exponential linear unit (SELU) as activation function with the standard scale $\lambda = 1.05070098$ and $\alpha = 1.67326324$ as they were calculated in (Klambauer et al. 2017).

$$selu(x) = \lambda \begin{cases} x & \text{if } x > 0 \\ \alpha e^x - \alpha & \text{if } x \leq 0 \end{cases} \quad (6)$$

7.3.1 Pure Convolutional Architecture




Figure 5: Encoder network architecture with an encoding size of 1024. The kernel size is 3×3 in every layer. The reparameterization trick in the last layer refers to the sampling as described in section 5.7.




Figure 6: Decoder network architecture with an encoding size of 1024. The kernel size is 3×3 in every layer. The last deconvolutional layer has a sigmoid activation function

7.3.2 Max Pooling Architecture

7.4 Understanding Latent Space

To understand and work with the latent space learned by the variational autoencoder is problematic and difficult since it is a very high dimensional space and it is not known what information of the input the autoencoder learns. For the understanding and visualization of the latent space its dimensionality needs to be reduced while preserving the underlying structure of the latent space. Then the low dimensional codes can be visualized in a way that allows humans to comprehend what the autoencoder learned. One dimensionality reduction method that could be used here is principal component analysis (PCA).

However, t-distributed stochastic neighbor embedding (t-SNE), as presented in 5.8, is the technique used in this work. The reason for this choice is that PCA focuses on preserving large distances, so dissimilar points in the high dimensional space appear far apart in a low dimensionality visualization. For visualization though it may be more interesting to keep more of the underlying structure which is accomplished by t-SNE as its focus is on preserving the context of points to their neighbors (Maaten et al. 2008). That behavior is caused by the asymmetry of the used Kullback-Leibler divergence 5.6. In the training process it penalizes a lot if large probabilities, i. e. far apart points, in high dimensional space are represented by small probabilities, i. e. close together points, in low dimensional space. Meanwhile, the penalty is little if small probabilities in high dimensional space are represented by large probabilities in low dimensional space, leading to the claimed focus on local similarities.

After t-SNE is used to reduce the dimensions of the latent code to 2, the low dimensional representation can be plotted in a scatter plot. That way it becomes evident if the autoencoder has learned any distinct clusters. To recap, every point in those plots is a low dimensional representation of a latent code that corresponds to a 128×128 section of an RGB satellite image. Now the goal is to find out what features the learned clusters in the plots correspond to. For that

purpose the semantic labeling and digital surface models 7.2 are used. In the first type of plot the points are colored according to the most dominant class in the corresponding image. Those classes are building, vegetation, water, ground and clutter. In the second kind of plot the points have a gray value equal to the average pixel values in the digital surface models meaning that points for images, that have on average more heights, are brighter than points representing images with less heights. In these two types of plots the points are squares if the related image is from Jacksonville or circles if it is from Omaha. In a third type of plot the points are represented by the corresponding images themselves. In those plots it can be observed if the clusters correlate to one of the visualized features which would mean that the autoencoder has learned to cluster that feature. This method is repeated for different coding sizes for the two architectures presented in 7.3. That process gives insight into what coding sizes yield the best unsupervised clustering. The specific experiments and results are found in the experiments section 8.2.

8 Experiments

8.1 Variational Autoencoder Architecture

8.2 Latent Space

9 Conclusion and Future Work

9.1 Conclusion

9.2 Future Work

References

- Bosch, Marc et al. (2019). “Semantic Stereo for Incidental Satellite Images”. In: *2019 IEEE Winter Conference on Applications of Computer Vision (WACV)*. IEEE, pp. 1524–1532.
- Chervinskii (2015). *Schematic picture of an autoencoder architecture*. [Online; accessed August 09, 2019][CC BY-SA 3.0 (<http://creativecommons.org/licenses/by-sa/3.0/>)].
- Chrislb (2005). *Diagram of an artificial neuron*. [Online; accessed August 08, 2019][CC BY-SA 3.0 (<http://creativecommons.org/licenses/by-sa/3.0/>)].
- Doersch, Carl (2016). “Tutorial on variational autoencoders”. In: *arXiv preprint arXiv:1606.05908*.
- Géron, Aurélien (2017). *Hands-on machine learning with Scikit-Learn and TensorFlow : concepts, tools, and techniques to build intelligent systems*. Sebastopol, CA: O’Reilly Media. ISBN: 978-1491962299.
- Glosser.ca (2013). *Artificial neural network with layer coloring*. [Online; accessed August 08, 2019][CC BY-SA 3.0 (<http://creativecommons.org/licenses/by-sa/3.0/>)].
- Goodfellow, Ian et al. (2016). *Deep learning*. MIT press.
- Klambauer, Günter et al. (2017). “Self-Normalizing Neural Networks”. In: *CoRR* abs/1706.02515.
- Krizhevsky, Alex et al. (2012). “Imagenet classification with deep convolutional neural networks”. In: *Advances in neural information processing systems*, pp. 1097–1105.
- Maaten, Laurens van der et al. (2008). “Visualizing Data using t-SNE”. In: *Journal of Machine Learning Research* 9, pp. 2579–2605.
- Martín Abadi et al. (2015). *TensorFlow: Large-Scale Machine Learning on Heterogeneous Systems*. Software available from tensorflow.org.
- Mishkin, Dmytro et al. (2016). “Systematic evaluation of CNN advances on the ImageNet”. In: *CoRR* abs/1606.02228.
- Pedregosa, F. et al. (2011). “Scikit-learn: Machine Learning in Python”. In: *Journal of Machine Learning Research* 12, pp. 2825–2830.
- Rossum, Guido (1995). *Python Reference Manual*. Tech. rep. Amsterdam, The Netherlands, The Netherlands.
- Springenberg, Jost Tobias et al. (2015). “Striving for Simplicity: The All Convolutional Net”. In: *ICLR (Workshop)*.
- Theis, Lucas et al. (Nov. 2015). “A note on the evaluation of generative models”. In: